

# Exercise

## Exercise Dataset

### Data Overview

The “AI-Powered Job Market Insights” dataset provides a snapshot of the modern job market, particularly focusing on the role of artificial intelligence (AI) and automation across various industries.

This dataset includes 500 unique job listings, each characterized by different factors like industry, company size, AI adoption level, automation risk, required skills, and job growth projections.

### Data Preview

```
ai_jobs |>
  gt_preview() |>
  tab_header(title = "AI-Powered Job Market Insights") |>
  tab_source_note(source_note = "Source: Kaggle https://www.kaggle.com/datasets/uom190346a/ai-powered-job-market-insights")
```

AI-Powered Job Market Insights

	job_title	industry	company_size	location	ai_adoption_level	automation_risk	required_skills	salary_remote	job_growth_projection
1	Cyber-security Analyst	Entertainment	Small	Dubai	Medium	High	UX/UI Design	111392.17	Yes Growth
2	Marketing Specialist	Technology	Large	Singapore	Medium	High	Marketing	73792.56	No Decline

Source: Kaggle <https://www.kaggle.com/datasets/uom190346a/ai-powered-job-market-insights>

## AI-Powered Job Market Insights

	job_title	industry	company_size	location	ai_adoption_level	automation_risk	required_skills	remote_friendly	job_growth_projection
3	AI Researcher	Technology	Large	Singapore	Medium	Low	UX/UI Design	Yes	Growth
4	Sales Manager	Retail	Small	Berlin	Low	Medium	Project Management	No	Growth
5	Cybersecurity Analyst	Entertainment	Small	Tokyo	Low	Low	Java Script	Yes	Decline
6..499									
500	HR Manager	Entertainment	Medium	Berlin	Medium	High	Project Management	Yes	Decline

Source: Kaggle <https://www.kaggle.com/datasets/uom190346a/ai-powered-job-market-insights>

To simplify our later code, I have created a separate table which is already filtered for the groups we will be looking at:

```
ai_jobs_risk <- ai_jobs |>
  filter(automation_risk %in% c("Low", "High"))

ai_jobs_high <- ai_jobs_risk |>
  filter(automation_risk == "High")
```

```
ai_jobs_low <- ai_jobs_risk |>
  filter(automation_risk == "Low")
```

### Dataset Features:

Source: Kaggle <https://www.kaggle.com/datasets/uom190346a/ai-powered-job-market-insights>

1. **Job\_Title:**
  - **Description:** The title of the job role.
  - **Type:** Categorical
  - **Example Values:** “Data Scientist”, “Software Engineer”, “HR Manager”
2. **Industry:**
  - **Description:** The industry in which the job is located.
  - **Type:** Categorical
  - **Example Values:** “Technology”, “Healthcare”, “Finance”
3. **Company\_Size:**
  - **Description:** The size of the company offering the job.
  - **Type:** Ordinal
  - **Categories:** “Small”, “Medium”, “Large”
4. **Location:**
  - **Description:** The geographic location of the job.
  - **Type:** Categorical
  - **Example Values:** “New York”, “San Francisco”, “London”
5. **AI\_Adoption\_Level:**
  - **Description:** The extent to which the company has adopted AI in its operations.
  - **Type:** Ordinal
  - **Categories:** “Low”, “Medium”, “High”
6. **Automation\_Risk:**
  - **Description:** The estimated risk that the job could be automated within the next 10 years.
  - **Type:** Ordinal
  - **Categories:** “Low”, “Medium”, “High”
7. **Required\_Skills:**
  - **Description:** The key skills required for the job role.
  - **Type:** Categorical
  - **Example Values:** “Python”, “Data Analysis”, “Project Management”
8. **Salary\_USD:**
  - **Description:** The annual salary offered for the job in USD.
  - **Type:** Numerical
  - **Value Range:** \$30,000 - \$200,000
9. **Remote\_Friendly:**
  - **Description:** Indicates whether the job can be performed remotely.
  - **Type:** Categorical

- **Categories:** “Yes”, “No”

10. **Job\_Growth\_Projection:**

- **Description:** The projected growth or decline of the job role over the next five years.
- **Type:** Categorical
- **Categories:** “Decline”, “Stable”, “Growth”

## Exercise - Mystery bags

To begin, split into three groups. Decide on the following roles, one per person:

1. **Sampler** - draws samples from the population
2. **Recording** - records the draws in R

Your group has been presented with a population model (bag containing salaries). There may be multiple populations represented in the various bags around the room.

**Please DO NOT look in the bag OR empty the contents of the bag.**

## Exercise - Mystery bags {nonincremental}

### The task

1. The **Sampler** draws 10 slips from the bag - this is your sample of size 5.
2. The **Recorder** records the values on the slips into the `sampling-exercise.R` file.
  - Create a new list for each sample (`sample_1 <- c(1, 2, 3, 2, 2)`)
3. Calculate the mean for the new sample and add it to the `sample_means` table.
4. Run the code to plot the histogram of your samples.
5. Return the slips to the bag.
6. Ensure the data is well mixed between samples.
7. Repeat

## Bibliography